

CSE 252B: Computer Vision II, Winter 2019 – Assignment 1

Instructor: Ben Ochoa

Due: Wednesday, January 16, 2019, 11:59 PM

Instructions

- Review the academic integrity and collaboration policies on the course website.
- This assignment must be completed individually.
- This assignment contains both math and programming problems.
- All solutions must be written in this notebook
- Math problems must be done in Markdown/LATEX. Remember to show work and describe your solution.
- Programming aspects of this assignment must be completed using Python in this notebook.
- This notebook contains skeleton code, which should not be modified (This is important for standardization to facilitate efficient grading).
- You may use python packages for basic linear algebra, but you may not use packages that directly solve the problem. Ask the instructor if in doubt.
- You must submit this notebook exported as a pdf. You must also submit this notebook as an .ipynb file.
- You must submit both files (.pdf and .ipynb) on Gradescope. You must mark each problem on Gradescope in the pdf.
- It is highly recommended that you begin working on this assignment early.

Problem 1 (Programming): Feature detection (20 points)

Download input data from the course website. The file price_center20.JPG contains image 1 and the file price_center21.JPG contains image 2.

For each input image, calculate an image where each pixel value is the minor eigenvalue of the gradient matrix

$$N = \begin{bmatrix} \sum_w I_x^2 & \sum_w I_x I_y \\ \sum_w I_x I_y & \sum_w I_y^2 \end{bmatrix}$$

where w is the window about the pixel, and I_x and I_y are the gradient images in the x and y direction, respectively. Calculate the gradient images using the fivepoint central difference operator. Set resulting values that are below a specified threshold value to zero (hint: calculating the mean instead of the sum in N allows for adjusting the size of the window without changing the threshold value). Apply an operation that suppresses (sets to 0) local (i.e., about a window) nonmaximum pixel values in the minor eigenvalue image. Vary these parameters such that around 600–650 features are detected in each image. For resulting nonzero pixel values, determine the subpixel feature coordinate using the Forstner corner point operator.

Report your final values for:

- the size of the feature detection window (i.e. the size of the window used to calculate the elements in the gradient matrix N)
- the minor eigenvalue threshold value
- the size of the local nonmaximum suppression window
- the resulting number of features detected (i.e. corners) in each image.



Display figures for:

- gradient image heat map before thresholding
- gradient image heat map after thresholding
- original image with detected features

My implementation takes around 25 seconds to run. If yours is more than 250 seconds you may lose points.

In [2]:

```
1 %matplotlib inline
2 import numpy as np
3 from PIL import Image
4 import matplotlib.pyplot as plt
5 import matplotlib.patches as patches
6 import time
7
8
9 # open the input images
10 I1 = np.array(Image.open('price_center20.JPG'), dtype='float')/255.
11 I2 = np.array(Image.open('price_center21.JPG'), dtype='float')/255.
12
13 # Display the input images
14 plt.figure(figsize=(14,8))
15 plt.subplot(1,2,1)
16 plt.imshow(I1)
17 plt.subplot(1,2,2)
18 plt.imshow(I2)
19 plt.show()
```



In [3]:

```
1 def rgb2gray(rgb):
2     return np.dot(rgb[...,:3], [0.299, 0.587, 0.114])
```

In [4]:

```
1 def AddandAverage(i, j, width, I):
2     return np.average(I[i - width : i + width + 1, j - width : j + width + 1])
```

In [5]:

```
1  def ImageGradient(I, w, t):
2      # inputs:
3      # I is the input image (may be mxn for Grayscale or mxnx3 for RGB)
4      # w is the size of the window used to compute the gradient matrix N
5      # t is the minor eigenvalue threshold
6      #
7      # outputs:
8      # N is the 2x2mxn gradient matrix
9      # b in the 2x1mxn vector used in the Forstner corner detector
10     # J0 is the mxn minor eigenvalue image of N before thresholding
11     # J1 is the mxn minor eigenvalue image of N after thresholding
12
13     m,n = I.shape[:2]
14     N = np.zeros((2,2,m,n))
15     b = np.zeros((2,1,m,n))
16     J0 = np.zeros((m,n))
17     J1 = np.zeros((m,n))
18
19     """your code here"""
20     #####Compute Gradient
21     Ix2, Iy2, IxIy = np.zeros(I.shape), np.zeros(I.shape), np.zeros(I.shape)
22     Ixb2, Iyb2 = np.zeros(I.shape), np.zeros(I.shape)
23     gradFilt = np.array([-1, 8, 0, -8, 1]) / 12
24
25
26     for i in range(2, I.shape[1] - 2): ##x
27         for j in range(2, I.shape[0] - 2): ##y
28             ix = np.dot(I[j, i - 2 : i + 3], gradFilt)
29             iy = np.dot(I[j - 2 : j + 3, i], gradFilt.T)
30             Ix2[j, i], Iy2[j, i], IxIy[j, i] = ix **2, iy **2, ix * iy
31             Ixb2[j, i], Iyb2[j, i] = Ix2[j, i] * i + IxIy[j, i] * j, Iy2[j, i]
32
33     width = (w - 1) // 2
34     for i in range(width, m - width):
35         for j in range(width, n - width):
36             ix2, iy2, ixiy = AddandAverage(i, j, width, Ix2), AddandAverage(i,
37             Ixb2, Iyb2 = AddandAverage(i, j, width, Ixb2), AddandAverage(i, j,
38             N[:, :, i, j] = np.matrix([[ix2, ixiy], [ixiy, iy2]])
39             b[:, :, i, j] = np.matrix([[Ixb2], [Iyb2]])
40             w, v = np.linalg.eig(N[:, :, i, j])
41             J0[i, j] = np.min(w)
42             if J0[i, j] > t:
43                 J1[i, j] = J0[i, j]
44
45     return N, b, J0, J1
46
47 def NMS(J, w_nms):
48     # Apply nonmaximum supression to J using window w
49     # For any window in J, the result should only contain 1 nonzero value
50     # In the case of multiple identical maxima in the same window,
51     # the tie may be broken arbitrarily
52     #
```

```

53 # inputs:
54 # J is the minor eigenvalue image input image after thresholding
55 # w_nms is the size of the local nonmaximum suppression window
56 #
57 # outputs:
58 # J2 is the mxn resulting image after applying nonmaximum suppression
59 #
60
61 J2 = J.copy()
62 """your code here"""
63 Jmax = np.zeros(J2.shape)
64 width = (w_nms - 1) // 2
65 for i in range(width, J.shape[0]):
66     for j in range(width, J.shape[1]):
67         Jmax[i, j] = np.max(J2[i - width : i + width + 1, j - width : j + w
68
69 for i in range(width, J.shape[0]):
70     for j in range(width, J.shape[1]):
71         if J[i, j] < Jmax[i, j]:
72             J2[i, j] = 0
73         else:
74             J2[i, j] = Jmax[i, j]
75
76 return J2
77
78 def ForstnerCornerDetector(J, N, b):
79     # Gather the coordinates of the nonzero pixels in J
80     # Then compute the sub pixel location of each point using the Forstner oper
81     #
82     # inputs:
83     # J is the NMS image
84     # N is the 2x2mxn gradient matrix
85     # b is the 2x1mxn vector computed in the image_gradient function
86     #
87     # outputs:
88     # C is the number of corners detected in each image
89     # pts is the 2xC list of coordinates of subpixel accurate corners
90     # found using the Forstner corner detector
91
92     """your code here"""
93     pts = np.zeros((2, 1))
94     C = 0
95     for i in range(0, J.shape[0]):
96         for j in range(0, J.shape[1]):
97             if J[i, j] != 0:
98                 C += 1
99                 pts = np.hstack((pts, np.dot(np.linalg.inv(N[:, :, i, j]), b[:,
100
101     return C, pts[:, 1:]
102
103
104 # feature detection
105 def RunFeatureDetection(I, w, t, w_nms):
106     N, b, J0, J1 = ImageGradient(I, w, t)

```

```

107     J2 = NMS(J1, w_nms)
108     C, pts = ForstnerCornerDetector(J2, N, b)
109     return C, pts, J0, J1, J2
110

```

In [145]:

```

1  # input images
2  I1 = np.array(Image.open('price_center20.JPG'), dtype='float')/255.
3  I2 = np.array(Image.open('price_center21.JPG'), dtype='float')/255.
4
5  # parameters to tune
6  w = 7
7  t = 0.0007
8  w_nms = 7
9
10 tic = time.time()
11 # run feature detection algorithm on input images
12 C1, pts1, J1_0, J1_1, J1_2 = RunFeatureDetection(rgb2gray(I1), w, t, w_nms)
13 C2, pts2, J2_0, J2_1, J2_2 = RunFeatureDetection(rgb2gray(I2), w, t, w_nms)
14 toc = time.time() - tic
15
16 print('took %f secs'%toc)
17
18 # display results
19 plt.figure(figsize=(14,24))
20
21 # show pre-thresholded corner heat map from gradient image function
22 plt.subplot(3,2,1)
23 plt.imshow(J1_0, cmap='gray')
24 plt.title('pre-thresholded gradient image')
25 plt.subplot(3,2,2)
26 plt.imshow(J2_0, cmap='gray')
27 plt.title('pre-thresholded gradient image')
28
29 # show thresholded corner heat map from gradient image function
30 plt.subplot(3,2,3)
31 plt.imshow(J1_1, cmap='gray')
32 plt.title('thresholded gradient image')
33 plt.subplot(3,2,4)
34 plt.imshow(J2_1, cmap='gray')
35 plt.title('thresholded gradient image')
36
37 # show corners on original images
38 ax = plt.subplot(3,2,5)
39 plt.imshow(I1)
40 for i in range(C1): # draw rectangles of size w around corners
41     x,y = pts1[:,i]
42     ax.add_patch(patches.Rectangle((x-w/2,y-w/2),w,w, fill=False))
43 # plt.plot(pts1[0,:], pts1[1,:], '.b') # display subpixel corners
44 plt.title('found %d corners'%C1)
45

```


```

46 ax = plt.subplot(3,2,6)
47 plt.imshow(I2)
48 for i in range(C2):
49     x,y = pts2[:,i]
50     ax.add_patch(patches.Rectangle((x-w/2,y-w/2),w,w, fill=False))
51 # plt.plot(pts2[0,:], pts2[1,:], '.b')
52 plt.title('found %d corners'%C2)
53
54 plt.show()


```

took 70.646483 secs


pre-thresholded gradient image




pre-thresholded gradient image



thresholded gradient image




thresholded gradient image

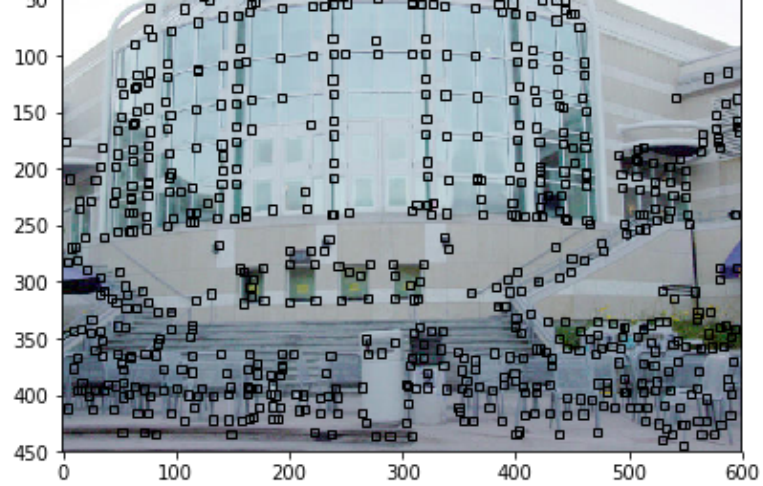



found 628 corners



found 650 corners





Final values for parameters

- $w = 7$
- $t = 0.0007$
- $w_{nms} = 7$
- $C1 = 628$
- $C2 = 650$

Problem 2 (Programming): Feature matching (15 points)

Determine the set of one-to-one putative feature correspondences by performing a brute-force search for the greatest correlation coefficient value (in the range $[-1, 1]$) between the detected features in image 1 and the detected features in image 2. Only allow matches that are above a specified correlation coefficient threshold value (note that calculating the correlation coefficient allows for adjusting the size of the matching window without changing the threshold value). Further, only allow matches that are above a specified distance ratio threshold value, where distance is measured to the next best match for a given feature. Vary these parameters such that around 200 putative feature correspondences are established. Optional: constrain the search to coordinates in image 2 that are within a proximity of the detected feature coordinates in image 1.

Report your final values for:

- the size of the matching window
- the correlation coefficient threshold
- the distance ratio threshold
- the size of the proximity window (if used)
- the resulting number of putative feature correspondences (i.e. matched features)

Display figures for:

- pair of images, where the matched features in each of the images are indicated by a square window about the feature

My implementation takes around 40 seconds to run. If yours is more than 400 seconds you may lose points.

In [155]:

```
1 import numpy.ma as ma
2 from math import sqrt
3
4 def NCC(I1, I2, pts1, pts2, w):
5     # compute the normalized cross correlation between image patches I1, I2
6     # result should be in the range [-1,1]
7     #
8     # inputs:
9     # I1, I2 are the input images
10    # pts1, pts2 are the point to be matched
11    # w is the size of the matching window to compute correlation coefficients
12    #
13    # output:
14    # normalized cross correlation matrix of scores between all windows in
15    # image 1 and all windows in image 2
16
17    """your code here"""
18    scores = np.zeros((pts1.shape[1], pts2.shape[1]))
19    R = (w - 1) // 2
20
21    pts1 = pts1.astype(int)
22    pts2 = pts2.astype(int)
23
24    for i in range(0, scores.shape[0]):
25        if pts1[1, i] < R or pts1[1, i] >= I1.shape[0] - R or pts1[0, i] < R or
26            continue
27        for j in range(0, scores.shape[1]):
28            if pts2[1, j] < R or pts2[1, j] >= I2.shape[0] - R or pts2[0, j] < R or
29                continue
30            w1 = I1[pts1[1, i] - R : pts1[1, i] + R + 1, pts1[0, i] - R : pts1[0, i] + R + 1]
31            w2 = I2[pts2[1, j] - R : pts2[1, j] + R + 1, pts2[0, j] - R : pts2[0, j] + R + 1]
32            mean1, mean2 = np.mean(w1), np.mean(w2)
33            matching_score = 0
34            denom1, denom2 = 0, 0
35            for r in range(0, w):
36                for c in range(0, w):
37                    matching_score += (w1[r, c] - mean1) * (w2[r, c] - mean2)
38                    denom1, denom2 = denom1 + (w1[r, c] - mean1) ** 2, denom2 + (w2[r, c] - mean2) ** 2
39            scores[i, j] = (matching_score / sqrt(denom1 * denom2))
40
41    return scores
42
43
44 def Match(scores, t, d, p):
45     # perform the one-to-one correspondence matching on the correlation coefficients
46     #
47     # inputs:
48     # scores is the NCC matrix
49     # t is the correlation coefficient threshold
50     # d distance ration threshold
51     # p is the size of the proximity window
```

```

52 #
53 # output:
54 # list of the feature coordinates in image 1 and image 2
55
56 """your code here"""
57 #inds = np.vstack((np.random.choice(pts1.shape[1],200,replace=False),
58 #                    np.random.choice(pts1.shape[1],200,replace=False)))
59
60 mask = ma.array(scores)
61 inds = np.zeros((2, 1))
62
63 while np.max(mask) > t:
64     maximum = np.max(mask)
65     index = np.unravel_index(np.argmax(mask), scores.shape)
66     #print(maximum, index, np.argmax(mask), scores.shape)
67     mask[index] = ma.masked
68     nextbest = max(mask[index[0],:].max(), mask[:, index[1]].max())
69     if (1 - maximum) < ((1 - nextbest) * d):
70         inds = np.hstack((inds, np.array([[index[0]], [index[1]]])))
71     mask[index[0], :] = ma.masked
72     mask[:, index[1]] = ma.masked
73
74 #print(mask)
75 return inds.astype(int)
76
77
78
79 def RunFeatureMatching(I1, I2, pts1, pts2, w, t, d, p):
80     # inputs:
81     # I1, I2 are the input images
82     # pts1, pts2 are the point to be matched
83     # w is the size of the matching window to compute correlation coefficients
84     # t is the correlation coefficient threshold
85     # d distance ration threshold
86     # p is the size of the proximity window
87     #
88     # outputs:
89     # inds is a 2xk matrix of matches where inds[0,i] indexs a point pts1
90     #     and inds[1,i] indexs a point in pts2, where k is the number of matches
91
92     scores = NCC(I1, I2, pts1, pts2, w)
93     inds = Match(scores, t, d, p)
94     return inds[:, 1:]
95

```

In [156]:

```

1 # parameters to tune
2 w = 7
3 t = 0.8
4 d = 0.8
5 p = np.inf

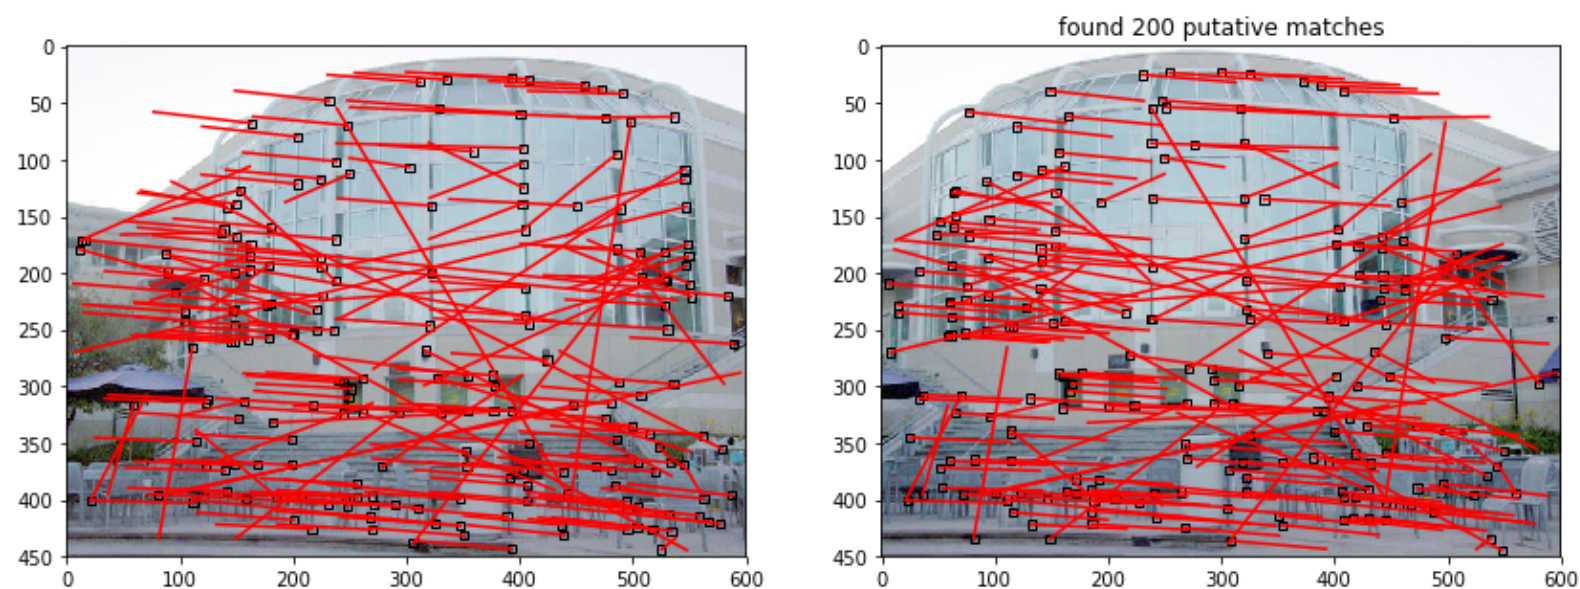
```

```

6
7 tic = time.time()
8 # run the feature matching algorithm on the input images and detected features
9 inds = RunFeatureMatching(rgb2gray(I1), rgb2gray(I2), pts1, pts2, w, t, d, p)
10 toc = time.time() - tic
11
12 print('took %f secs'%toc)
13
14 # create new matrices of points which contain only the matched features
15 match1 = pts1[:,inds[0,:]]
16 match2 = pts2[:,inds[1,:]]
17
18 # display the results
19 plt.figure(figsize=(14,24))
20 ax1 = plt.subplot(1,2,1)
21 ax2 = plt.subplot(1,2,2)
22 ax1.imshow(I1)
23 ax2.imshow(I2)
24 plt.title('found %d putative matches'%match1.shape[1])
25 for i in range(match1.shape[1]):
26     x1,y1 = match1[:,i]
27     x2,y2 = match2[:,i]
28     ax1.plot([x1, x2],[y1, y2],'-r')
29     ax1.add_patch(patches.Rectangle((x1-w/2,y1-w/2),w,w, fill=False))
30     ax2.plot([x2, x1],[y2, y1],'-r')
31     ax2.add_patch(patches.Rectangle((x2-w/2,y2-w/2),w,w, fill=False))
32
33 plt.show()
34
35 # test 1-1
36 print('unique points in image 1: %d'%np.unique(inds[0,:]).shape[0])
37 print('unique points in image 2: %d'%np.unique(inds[1,:]).shape[0])

```

took 38.231385 secs



unique points in image 1: 200
unique points in image 2: 200

Final values for parameters

- $w = 7$
- $t = 0.8$
- $d = 0.8$
- $p = \text{not used}$
- $\text{num_matches} = 200$

In []:

1	
---	--